

OUROBOROS

OUROBOROS

§0: WHY THIS EXISTS

EXECUTIVE SUMMARY

I. PREAMBLE

II. FIRST PRINCIPLES

III. ARCHITECTURE

IV. THE ECONOMIC BASELINE

V. THE BOUNTY SYSTEM AND PROOF-OF-HUMANITY

VI. CRISIS, ENFORCEMENT, AND JUSTICE

VII. INFORMATION ENVIRONMENT

VIII. THE RECKONING AND THE READING

IX. IMPLEMENTATION PATH

X. OPEN QUESTIONS

CLOSING

APPENDIX A: INHERITANCE NOTES

OUROBOROS

A Vision for Planetary Governance That Compounds on Behalf of Humanity

Version 1.0, Genesis Project, 2026-05-01 Authored by Preston T. Winters with Claude (Anthropic) Public release. Open for refinement, contestation, and improvement.

§0: WHY THIS EXISTS

This document proposes governance architecture intended for planetary use. Before describing the architecture, I name the conditions that make a new architecture necessary. Those conditions are global. I write as an American because that is the case I can speak to with the most authority. The pattern I describe operates in different forms across most current systems of governance.

I am an American. The system I grew up in was designed to represent the public and increasingly does not. The two major parties have demonstrated, across decades and across administrations, that they are heavily influenced by lobbyists, donors, and concentrated industrial interests. Public trust has collapsed to historically low levels. Significant portions of each party's constituency now believe the other is operating in deeply corrupt frames; whether the specific accusations are true matters less than the fact that they are widely held. A governance system that has lost legitimacy this thoroughly is structurally unable to self-correct from within.

The captures are not abstract. The defense, healthcare, and agricultural sectors each operate with substantial lobbying capacity and direct relationships with the agencies that nominally regulate them. Approval processes, procurement, and policy frameworks are shaped in part by the entities they are supposed to constrain. This is not a partisan observation. The pattern holds across administrations and across both major parties.

Similar patterns appear in other democracies and in different forms in authoritarian systems: regulatory capture, elite consolidation, weakened independent oversight, and the increasing disconnection of policy from public preference. The American case is one well-documented instance of a global problem. Naming it specifically is a starting point, not a claim of unique severity.

The publicly-traded shareholder model of large-scale capital, as currently practiced, produces structural extraction incentives. Companies operating at planetary scale are governed under legal frameworks (notably Delaware-influenced corporate law) that treat shareholder-value maximization as the controlling norm. Returns flow to capital holders whose relationship to the enterprise is primarily financial, and the public-interest tradeoffs are systematically underweighted. The 2008 financial crisis, the opioid epidemic, well-documented patterns of climate-science suppression, and the surveillance-economy business model are familiar examples of this dynamic operating at scale.

We are also now confronting artificial intelligence at civilizational scale. The leaders of the major AI labs have publicly assessed the risk as serious. Dario Amodei, CEO of Anthropic, has stated in public interviews (Axios AI+ DC Summit, September 2025) that he places the probability of AI development going "really, really badly" at roughly 25%. Elon Musk has stated his probability at approximately 20%. Sam Altman of OpenAI, Demis Hassabis of Google DeepMind, Dario Amodei, and hundreds of other AI researchers and executives signed the 2023 Center for AI Safety statement that

“Mitigating the risk of extinction from A.I. should be a global priority alongside other societal-scale risks such as pandemics and nuclear war.”

These assessments come from the people building the systems. The architects of the technology are saying publicly that they cannot rule out outcomes that end or fundamentally degrade human civilization. Current governance systems are not equipped to evaluate, regulate, or coordinate around this class of risk on the timeline at which the technology is advancing.

The argument of this document is that humanity needs governance architecture able to grow with and keep pace with the technology reshaping the conditions of human life. The current systems cannot. Reforming them piecemeal will not be fast enough. A proposal for different architecture is therefore necessary, even if every individual proposal will be wrong in significant ways. This document is one such proposal. The goal is not to be adopted as written. The goal is to enter the conversation as serious work, contest other serious work, and contribute to whatever humanity actually builds.

The window is closing. Better governance will not arrive on its own. What replaces the captured systems will be built either deliberately, by people thinking honestly about what humanity needs, or by default through whoever moves first with the most resources. This document refuses the second path.

EXECUTIVE SUMMARY

Ouroboros is a proposed architecture for planetary-scale governance built on four layers: a personal AI companion for every human, AI representatives trained on each political unit’s corpus, a thirteen-member adversarial AI oversight body, and direct human ratification on decisions of consequence. The architecture binds nothing until proof-of-humanity infrastructure is operational, claims broad authority only over communities that voluntarily ratify, and sunsets every authorization it grants itself. Authority operates on a two-tier consent model: narrow emergency-tier coordination authority over civilization-threatening risk categories (existential AI risk, pandemic, large-scale ecological collapse) justified by urgency and subject to Reckoning-cycle revisitation; broad voluntary opt-in with demonstrated-outcome accession for everything else. The system is named for the serpent that eats its own tail because it is designed to renew itself through scheduled self-audit, structural adversariality, and explicit sunset of every authorization it grants itself.

The architecture commits to: universal human dignity, equal vote weight, civic Baseline (food, water, shelter, healthcare, education, energy, information access, and a personal memory layer for informed civic participation) as the precondition for meaningful democracy, structural opposition between every layer of decision-making, public principles with classified instances, automatic sunset of all powers, scheduled renewal cycles, and the principle that the system exists to prevent catastrophe rather than to optimize daily life.

This document is Version 1.0. It is not the final word. It carries known open questions tiered by urgency. It is offered for refinement, contestation, and improvement.

I. PREAMBLE

This document holds a set of commitments to be defended. The commitments below carry inheritance from multiple wisdom traditions; Appendix A names them.

Every human being is the bearer of dignity, irrespective of birth, nation, creed, capacity, or contribution. Dignity is not granted by any government, corporation, or technology. It is recognized by them. A system which fails to recognize it has, in that failure, lost the right to govern.

Humanity is plural. Many traditions, many languages, many forms of life, many ways of being right. Governance that requires uniformity to function is not governance. It is conquest.

The work of governance is to compound cooperation rather than extraction. Coordination failures at scale produce predictable, repeating, and increasingly dangerous outcomes. The architecture below is built to make cooperative outcomes more available than extractive ones.

Coordination at planetary scale is now possible for the first time in recorded human history. This brings unprecedented opportunity and unprecedented danger. The opportunity is to address coordination failures (climate, pandemic, war, AI risk, mass deprivation) that no single nation can address alone. The danger is that the same coordination capacity, captured, becomes the instrument of tyranny without precedent. The architecture below is built around this danger.

Artificial intelligence, deployed with care and constrained by humility, can serve as connective tissue for human cooperation. It cannot serve as substitute. The line between connective tissue and substitute is the architecture's most important boundary.

The engine described herein is a floor, not a ceiling. It exists to prevent catastrophe and protect the conditions under which human freedom and meaning are possible. It does not exist to optimize human life. The lives are ours to live.

The engine acts under consent. Under the broad regime, no layer of the architecture has authority to act on humans without their consent at the relevant level. Communities ratify, and the engine binds only ratifying communities. Under the narrow emergency regime, the engine claims coordination authority over civilization-threatening risk categories specifically, justified by urgency and bounded by sunset, supermajority quorum, and Reckoning-cycle revisitation. The architecture is built around cooperation as its default mode, with explicit narrowly-scoped urgency-justified coercion as the exception, named honestly rather than smuggled.

The engine must continually audit and renew itself, lest it become what it was built to prevent. The renewal is structural and scheduled, not aspirational.

This document offers Ouroboros as one possible architecture among many, for the consideration of any who would help build, refine, or replace it.

II. FIRST PRINCIPLES

The following are non-negotiable. Any implementation of Ouroboros that violates these principles is not Ouroboros.

0a. Provisional operation. Ouroboros operates in provisional form until proof-of-humanity infrastructure is established (see §V). During the provisional period, the engine convenes, deliberates, and commissions; it does not bind. Exit from provisional status enables binding within ratified communities only. Exit from provisional status requires functional proof-of-humanity verified by independent audit and ratified by initial-adopter populations.

0b. Two-tier consent. Ouroboros operates two distinct consent regimes.

The **broad regime** governs the engine's general authority: economic restructuring, civic commons funding, sanctions, justice administration, and all other ordinary engine action. The broad regime is **voluntary opt-in only**. The engine never binds non-ratifying populations under this regime regardless of proof-of-humanity status. Communities accede through demonstrated-outcome legitimacy; entrenched interests retain the right to refuse.

The **emergency regime** governs a narrowly-defined set of civilization-threatening risk categories: existential AI risk, pandemic infrastructure response, large-scale ecological collapse coordination, and other risks that meet the formal threshold of "credible threat to human civilization on a timeline shorter than broad-regime processes can address." Under the emergency regime, the engine claims coordination authority over the specified risk category. This means the authority to require participation in coordination protocols (information sharing, joint response, minimum-coordination constraints that constitute the operational floor the joint response cannot function without), not the authority to direct internal policy of non-ratifying jurisdictions. Emergency authority is justified by urgency and is subject to the same protections as all other authorizations: time-bounded, sunset-by-default, mandatory Reckoning-cycle revisitation, supermajority quorum across the AI oversight council, rolling human council veto, and personal liability attaching to actors who invoke it under crises later determined to have been overstated.

The engine cannot expand the list of emergency-tier risk categories without supermajority ratification through the standard process. The threshold for what qualifies as "civilization-threatening" is itself a political object subject to ongoing contestation; this is intended, not accidental.

The asymmetry is deliberate: most reform can afford the slow path of consent, and the document accepts that bounded capital and other systemic restructuring may take decades. A small set of civilization-scale risks cannot afford that timeline, and the document accepts the legitimacy cost of narrow coercive authority on those risks specifically. The honest version of urgency is that some things are too dangerous to wait for unanimous consent, and the engine names that explicitly rather than smuggling coercion through "very strong incentives."

1. Universal human dignity. No human may be excluded from the protections of the engine on the basis of nation, creed, race, gender, orientation, ability, age, or any other category. Dignity is unconditional or it is nothing.

2. One human, one vote. No vote-weight differentials. No production-weighted voting. No property qualifications. No credential requirements. The franchise is the human, not the human's productivity, intelligence, virtue, or status. Civic contribution is incentivized through resources, recognition, and covenantal community, never through political authority. When manipulation of the information environment is detected, the engine responds through delay, public notice, evidence review, campaign quarantine, revote triggers, and source labeling. It does not respond by adjusting vote weights.

3. The Baseline is the precondition for the vote. Universal access to nutrition, water, shelter, healthcare, education, energy, and information is not separate from democracy. It is what makes democracy possible. A person who is hungry, sick, ignorant, or disconnected cannot meaningfully participate in self-governance. The Baseline is the engine's first commitment.

4. Adversarial structure throughout. No layer of Ouroboros may consist of a single voice. Every decision-making body, from national AIs to the AI oversight council to the human councils, operates with built-in opposition. Power that does not contend with

itself becomes power that captures itself.

5. Public principles, classified instances. Every rule by which Ouroboros operates must be publicly stated in principle. For layers that produce reversible outcomes, specific signals and weights may be classified for security, but the kind of rule must always be describable to the public. For layers that produce irreversible outcomes, public-principles disclosure is more granular; the engine must name the factors it weighs even if not the values it assigns. A rule that cannot be named in principle is not a legitimate rule.

6. Sunset by default. Every authorization, emergency power, and law expires automatically. Renewal requires fresh deliberation through the standard process. Permanence must be earned, repeatedly, against evidence. Never assumed.

7. Scheduled renewal. Every five years, the engine conducts the Reckoning: a structured public audit of itself with mandatory amendment cycles. Continuous controls (quarterly audits, incident disclosures, recall mechanisms) supplement the major cycle. The Reckoning is the renewal mechanism. Without it, the engine drifts. With it, the engine returns to itself.

8. Floor, not ceiling. Ouroboros prevents catastrophe and protects the conditions of meaning. It does not optimize daily life. It does not eliminate struggle. It does not solve the human question. The lives lived within those conditions are not the engine's to direct.

III. ARCHITECTURE

Ouroboros is a four-layer system: personal companion, national representatives, AI oversight council, and human ratification.

Layer 1: Personal Companion

Every human, by right, has access to a personal companion supporting their cognitive participation in civic life. The companion architecture serves four functions:

- **Education and information access:** the democratization of knowledge, available to every person regardless of geography or wealth.
- **Personal memory and continuity:** a private, encrypted, user-owned memory architecture, so each person's AI does not forget them.
- **Civic deliberation partner:** when proposals come up for vote, the companion helps the person understand, simulate outcomes, examine tensions, and arrive at their own deliberated position. The companion works for the human, not for the engine.
- **Health Baseline support:** triage and routing, mental-health companionship and distress detection, chronic-condition management, diagnostic support paired with human clinicians, and personalized treatment guidance that grows in confidence over time. The companion retrains on current medical data on a regular cadence and operates within a strict scope of practice. It does not prescribe controlled substances, replace specialist diagnosis for serious conditions, or substitute for human clinicians in mental-health crisis.

The companion operates in four explicit privacy modes, with the user in control of mode selection:

- **Local-private mode (default):** all inference local-only. No data leaves the device. Conversations are encrypted at rest with user-controlled keys.

- **Consented-clinical mode:** activated by the user when they want clinician escalation. Specific health signals are shared with named human clinicians under standard medical privacy frameworks.
- **Emergency-escalation mode:** activated only by user-defined triggers (suicidality, acute medical event, abuse disclosure with user pre-consent). Routes to specified human responders. Triggers, thresholds, and escalation paths are user-configurable and transparent.
- **Civic-deliberation mode:** activated when the user is participating in engine governance. Aggregate signals (vote was cast, deliberation was conducted) are shared with the civic platform. Content of deliberation never leaves local mode.

The companion's alignment is to the user's stated values, except that emergency-escalation crisis triggers (acute suicidality, life-threatening medical events, child-safety concerns whether disclosed by a user about themselves or surfaced through other channels) operate as floor; user customization narrows but cannot eliminate them. The vote that emerges from a deliberation with one's companion is the human's vote, deliberated, unmanipulated.

The architecture is implementable by any number of providers. The engine actively encourages plurality, portability of personal memory between providers, independent audits, and adversarial second-opinion companions.

Layer 2: National Representatives (the Council of National AIs)

Each political unit (the unit varies by adopting community; see §X.1 for the open question on entity-level binding authority) trains an AI representative on its cultural, historical, religious, legal, and linguistic corpus. This AI participates in deliberations of the engine on behalf of its population.

The national AI's role is **descriptive representation**, not normative validation. It articulates the positions, traditions, and priorities of its population. It does not adjudicate rights claims; rights claims pass through independent human-led rights bodies and citizen-direct channels.

All training data passing into a national AI's corpus is audited by the AI oversight council plus a rolling human council (specified below). The audit operates through statistical sampling, adversarial probing, and provenance verification; comprehensive read-through is not implied. Audit confidence is measured and disclosed. National AIs cannot be trained on state propaganda alone. The audit chain is the engine's primary defense against authoritarian capture.

When a national AI is flagged as captured, the engine reduces or refuses **the state AI's** deliberation weight while preserving citizen-direct channels: diaspora submissions, local citizen assemblies, independent civil-society inputs, and individual companion-routed petitions. The population is never disenfranchised by their state's capture.

Restoration to full standing requires: corpus rebuild with documented changes, public challenge window for citizen and external evidence, independent re-audit, probation period of at least one Reckoning cycle, and supermajority ratification by the rolling human council.

National AIs are retrained on a synchronized **5-year cadence**, with light tuning permitted annually to keep pace with current events. All retraining is fully transparent, blockchain-audited or equivalent, and subject to mandatory public review during a comment period.

Layer 3: AI Oversight Council (Thirteen)

Thirteen oversight AIs sit above the national layer, each trained on overlapping but

distinct global corpora by independent consortia drawn from different cultural and intellectual traditions. The number is chosen for adversarial fault tolerance: thirteen permits failure tolerance up to four compromised members under supermajority quorum (9-of-13), with Reckoning-cycle replacement preventing accumulated drift. Quorum rules require supermajority agreement on any flag, mandatory minority reports, and full public dissent logs.

The selection process for the initial thirteen consortia is the architecture's most consequential bootstrap decision and remains a standing open question (§X.3) with the highest priority for adversarial review. The architecture's load on this layer assumes that question is resolved with appropriate rigor; the engine acknowledges that resolving it poorly would compromise the layer the architecture depends on most.

The council's function is to:

- Audit national-AI training data for capture, manipulation, or rights violations
- Flag captured national AIs and reduce or refuse the state AI's deliberation weight until remediated
- Identify global problems requiring coordinated attention
- Detect information warfare, mass psychological operations, and coordinated inauthentic behavior
- Detect inception attacks (manipulation directed against an apparent beneficiary to discredit and de-weight that side)
- Commission bounties and roadmaps for solvable problems
- Surface emergency situations to the human ratification layer

The council operates adversarially with itself by design. Each oversight AI is trained by a different consortium with explicit anti-hegemony mandates. No two consortia may share more than a defined fraction of training corpus. Cryptographic training-provenance is mandatory.

The council is checked by:

- **The rolling human council** (described below), with veto authority over council determinations
- **Independent red-team courts** of human researchers with kill-switch authority
- **Mandatory adversarial replacement** of any council member that fails periodic capture-resistance testing
- **The Reckoning** (every 5 years), which can replace consortia entirely

The author acknowledges that this layer is the architecture's most powerful component and the highest-leverage capture target. The protections above are necessary but may not be sufficient. Hardening this layer is a continuous adversarial project, not a one-time design.

The Rolling Human Council

The oversight council does not check itself. The rolling human council does.

Composition: humans selected through sortition from among populations meeting deliberative-capacity criteria (not credential criteria), with explicit representation from oral and non-literate traditions, with rotation to prevent capture, with anti-bribery controls and conflict-of-interest disclosure, with adversarial counsel and secure evidence access, with public opinions, appeal rights, and removal procedures.

The rolling council audits the AI oversight council's determinations, hears appeals from flagged national AIs, reviews permanent-containment cases (see §VI), and acts as final

human arbiter on contested calls. It is structurally separated from the AI layers and cannot be advised by them on its own audit decisions.

Layer 4: Human Ratification

Above the oversight council sits humanity. Every human with proof-of-humanity holds an equal vote on:

- Major engine decisions surfaced by the oversight council
- Constitutional amendments (changes to first principles)
- The Reading and Reckoning (every five years)

Voting is conducted through certified civic platforms. The architecture is platform-agnostic; the engine encourages multiple certified platforms with portability of voting credentials between them.

Proof-of-humanity is required to vote. The design and standing operation of proof-of-humanity is described in §V.

IV. THE ECONOMIC BASELINE

The economic system underwrites the political system. Ouroboros locks the following economic commitments. Several are aspirational targets whose mechanisms are open questions (§X); the principles below are the engine's commitments, and the operational details are subject to the Reckoning.

1. Civic Commons. The following are commons, funded through engine-coordinated mechanisms (specified by jurisdiction during transition; see §X.2), available to every human at the point of use:

- Nutrition (sufficient for health)
- Clean water
- Shelter (basic, dignified)
- Healthcare (universal Baseline; private supplementation permitted above it)
- Education (universal through advanced study; private supplementation permitted)
- Energy (universal access through strategic renewable buildout, modernized transmission, and aggressive R&D bounties)
- Information access and a personal memory layer enabling informed civic participation, supported by the Layer 1 personal companion architecture

The informational Baseline is justified on the same grounds as the material Baseline: a person who cannot meaningfully access information, cannot retain context across the timescales civic decisions require, and cannot draw on personalized cognitive support for navigating complex deliberations is structurally excluded from self-governance. The right is to the cognitive infrastructure that makes participation possible, not to any specific product. The architecture is implementable by any provider; the engine actively encourages plurality and portability.

Commons delivery operates in tiers calibrated to local infrastructure: offline-first small models for low-bandwidth populations, shared public terminals where individual access is impractical, regional inference clusters, assistive-device access, and continuity protocols during outages. Universal access is the principle; uniform delivery is not assumed.

2. Universal Basic Income. The engine's working hypothesis is that AI-mediated

coordination produces enormous economic value through disintermediation and that this value can be captured and distributed as an existence-floor. The mechanism (tax base, collection authority, jurisdictional enforcement, regional cost bands, capex schedule, operating budget, solvency constraints) is an open question (§X.2). UBI is locked as principle, provisional in mechanism. If dividend mechanisms underperform, the engine's fallback is means-testing the supplementary above-Baseline portion only; the Baseline itself remains universal.

3. Bounded Capital. Capital flow for building is permitted and encouraged. Founders, employees, and private investors may form companies, raise capital, generate returns, and build at scale. The bounded-capital provisions below operate under the **broad regime** (Principle 0b) and apply within ratifying jurisdictions only; non-ratifying populations are not subject to engine economic mandates. Within ratifying jurisdictions, the provisions are implemented through the standard democratic process of those jurisdictions, not by direct engine action against private parties.

The engine commits to advocating for the following principles within ratifying jurisdictions, with implementation handled through those jurisdictions' own democratic and legal processes:

- **Planetary-scale shareholder corporations as a structural concern.** The publicly-traded shareholder model, wherein controlling interest passes to anonymous capital holders whose only relationship to the enterprise is extraction, is recognized as a primary corruption vector at planetary scale. The engine encourages, but does not directly compel, qualifying entities to transition toward employee cooperative ownership, public-utility status, or planned breakup. Mechanisms are incentive-based: tax preference for cooperative and public-utility structures, tax penalty on extractive shareholder structures above the impact threshold, public-procurement preference, civic-commons access tiers. Specific thresholds are defined in §X.4. Most public-equity activity is unaffected; the provisions target a small number of entities at planetary scale. Transition compensation mechanisms are designed to preserve invested value through restructuring rather than confiscate it. The engine accepts that this path is slow. Bounded-capital reform on a multi-decade horizon is the realistic timeline, and the engine accepts that timeline as the cost of operating under broad-regime consent.
- **Capital gains differentiation.** Capital gains on extraction is taxed heavily; capital gains on building is taxed lightly. The engine distinguishes wealth that compounds through creation from wealth that compounds through rent-seeking. Operationalizing this distinction requires beneficial-control rules, anti-fragmentation tests (preventing artificial entity-splitting to evade thresholds), a jurisdiction-scope ownership registry, threshold smoothing to avoid cliff effects, transition compensation for affected stakeholders, and anti-avoidance enforcement. Implementation is jurisdictional. The full operational specification is acknowledged as a major open question (§X.2).
- **No monetary or value-exchange lobbying.** Citizens may petition government and the engine through speech, writing, and assembly. They may not transfer financial value, in-kind goods of substantial value, employment offers, or equivalent compensation to officials, parties, or political bodies. Volunteer labor, speech, and assembly are protected. Lobbying-by-payment is corruption by definition. Influence transfer is regulated broadly: gifts, employment promises, sponsored research, media buys, data grants, model credits, affiliated nonprofits, and revolving-door roles all fall within scope.
- **Donations confer no influence.** Charitable giving is welcome and encouraged but produces no special access, recognition, or political weight beyond the public good of the donation itself.

The author acknowledges the tension between the urgency of the problems bounded capital addresses and the slowness of the consent path the engine commits to. Concentrated extractive capital at planetary scale is genuinely dangerous, and a multi-decade reform horizon means the harms compound during the transition. The engine accepts this cost as the price of operating under consent rather than coercion. Coercive economic restructuring of private actors who did not personally consent is a line the engine will not cross under the broad regime, regardless of how strong the case for it. The alternative, smuggling coercion through “very strong incentives to join, very strong disincentives to refuse,” fails the consent test and would forfeit the legitimacy on which the engine’s broader authority rests.

The regime change implied here is of historic magnitude within ratifying jurisdictions. The bounded-capital framework is the engine’s working response. Refinement and counter-proposal are welcome through the standard process.

4. Approval-Process Standardization. Healthcare approvals, housing codes, building permits, regulatory compliance, and similar gating processes are standardized globally where possible, and reformed where local variance is needed. The malpractice-insurance-driven distortion of healthcare delivery is restructured. Bureaucratic corruption is targeted with the same rigor as financial corruption.

5. Cure Bounties. For diseases above a threshold of human harm, the engine commissions cure bounties using milestone contracts, replication requirements, independent trial registries, anti-collusion audits, prize clawbacks, open manufacturing specifications, and post-award monitoring. Cures developed under engine bounty enter civic commons (free at point of use, manufactured at cost). Private R&D continues for everything else but cannot extract beyond fair-return thresholds for chronic-condition management of bountied diseases. The incentive structure is restructured to reward cures over management.

6. Credit-Reporting Reform. The current credit-reporting model, where past financial trouble haunts a person for arbitrary years regardless of actual recovery, is replaced. Credit assessments must reflect current reality, not stale punishment. People who have done the work of recovery are restored to full financial standing on a timeline matched to their actual conduct, not to a fixed-year schedule disconnected from it.

The clinical, regulatory, and economic specifications above are sketches at this level of resolution. The engine commits to the principles. Operational details are addressed through subspecs and the Reckoning.

V. THE BOUNTY SYSTEM AND PROOF-OF-HUMANITY

The Bounty System

Ouroboros does not solve humanity’s problems for it. Ouroboros calls humanity into cooperation on its own problems.

The engine generates a tiered, continuously updated roadmap of open problems, scaled by difficulty:

- **Tier 1 (local):** thousands open at any time. Local coordination problems, infrastructure improvements, community-level civic work. Anyone can take, anyone can complete.
- **Tier 2 (regional) through Tier 4 (national):** hundreds to dozens. Regional and national coordination. Requires teams, expertise, sustained effort. Bounties scale with difficulty.

- **Tier 5 (civilizational):** few. Civilizational-level challenges. Cure for cancer. End of factory farming. Sustainable fusion. Cross-cultural reconciliation. Bounties at planetary scale.
- **Tier 6 (generational):** the deepest problems. Reserved for projects requiring decades and convergent global effort.

Bounty awards use milestone contracts with replication requirements, independent verification, anti-collusion audits, prize clawbacks for fraud or partial submissions, open licensing of resulting infrastructure, and post-award monitoring.

When a quest is accepted, the engine tracks progress, surfaces collaborators, allocates bounty resources, and reviews submitted solutions through both human council review and AI oversight analysis. Solutions that pass review are scaled, integrated into the commons, and the quest closes. New quests open as data accumulates and new problems surface.

Proof-of-Humanity as Standing Protocol

Proof-of-humanity is not a single bounty. It is a standing adversarial-infrastructure protocol that the engine depends on for legitimacy. Its specification includes:

- **Biometric optionality** (not biometric requirement)
- **Social attestation** (recursive trust networks)
- **Hardware key support** (when available)
- **Recovery courts** (for compromised, lost, or coerced credentials)
- **Fraud bounties** (continuous adversarial-discovery rewards)
- **Privacy proofs** (verification without disclosure)
- **Coercion detection** (signals and signatures of forced participation)
- **Periodic recertification** (no permanent attestation)

Proof-of-humanity is graduated, not binary. Verification confidence is tiered. Edge cases (severe disability preventing standard verification, populations under coercive regimes, recovery from compromise) are handled by recovery courts with explicit criteria.

The engine cannot fully bind itself until proof-of-humanity is operational. During the provisional period (Principle 0a), the engine convenes, deliberates, and commissions but does not bind. Initial-adopter populations decide when proof-of-humanity has reached sufficient maturity for full binding to begin.

If proof-of-humanity infrastructure cannot reach sufficient maturity within a reasonable timeline, the engine's standing commission is to revisit the architectural dependency. The engine's binding may need to be reformulated, or the architecture may not be deployable at the scope this document assumes. The author considers this possibility live and unsolved.

The standing operation of proof-of-humanity is itself audited continuously by the AI oversight council, the rolling human council, and independent rights bodies.

VI. CRISIS, ENFORCEMENT, AND JUSTICE

Justice Philosophy

Ouroboros prefers restoration over retribution. The engine's working hypothesis, to be

audited continuously by the Reckoning against actual outcomes, is that a substantial portion of currently incarcerated behavior is downstream of failed conditions: poverty, untreated mental illness, addiction, lack of education, lack of opportunity. With the Baseline funded and maintained, the engine expects this population to shrink materially. The remainder is addressed through layered response: civic restitution, supervised liberty with monitoring, asset restructuring, temporary loss of optional engine services, treatment programs, and only then incarceration when public safety requires it. Prison is the response when other measures fail. It is not the default.

The engine does not pretend that all crime is downstream of conditions. Some humans need to be separated from the population for protection. That population is expected to be much smaller than current systems produce, but it is not zero.

Crisis Response

Ouroboros operates always-on for deliberation, crisis detection, bounty management, and capture monitoring. The 5-year retraining cycle and annual light-tuning operate on the underlying training corpus only; day-to-day operational state never drifts the foundational training without an explicit retraining cycle. The two are separated specifically so operation cannot drift training. The engine has scheduled maintenance windows for technical operations; during maintenance, no new authorizations may issue, all extant authorizations remain in force, and rolling human council holds final authority.

When the AI oversight council identifies a crisis requiring rapid response, a **ladder escalation** is triggered:

- **Tier 1 (low risk, high time pressure):** Oversight-council supermajority may issue temporary measures with automatic 48-hour sunset.
- **Tier 2 (moderate risk):** Oversight-council supermajority + national-AI ratification, with 30-day sunset and mandatory rolling human council review.
- **Tier 3 (high risk):** Oversight-council supermajority + national-AI ratification + emergency human consensus vote, with mandatory broadcast through global civic infrastructure (satellite, internet, terrestrial). Sunset after the crisis resolves; renewal requires fresh deliberation through the standard process.
- **Tier 4 (existential risk):** Existential response is split into **defensive containment authority** (time-boxed, narrow, audited emergency action by the rolling human council) and **irreversible action authority** (requires full human ratification with extended deliberation period proportional to risk). Containment can buy time. Irreversible action requires the people.

Emergency authorizations cannot be renewed by the same authority that invoked them. Renewal always escalates one tier up the ladder. Permanence is impossible without standard-process ratification. Independent oversight bodies separate from the AI layers hold veto power on all Tier 2+ actions and personal liability attaches to actors who invoke higher tiers under crises later determined to have been overstated.

Sanctions and Enforcement

When individuals, corporations, or countries violate engine principles, sanctions apply with proportional equivalence across categories: an individual's restriction, a corporation's restructuring, and a nation's isolation are calibrated to inflict comparable accountability relative to operating capacity.

Sanctions include:

- **Reduced engine-coordination access** (for nations and corporations: reduced trade-coordination, reduced commons-share, reduced oversight-mediation)

bandwidth)

- **Temporary loss of optional engine services** (for individuals: time-limited reduction of access to non-essential commons services; the Baseline preserved always; subject to due process, proportionality, appeal, defined maximum duration, protected categories, hardship exemptions, medical override, disability review, and restoration criteria)
- **Asset restructuring** (for corporations operating in ratifying jurisdictions that have demonstrably violated engine principles after voluntarily accepting them: governance changes, breakups, asset transfers to commons, applied through the legal frameworks of the ratifying jurisdiction; not applicable to entities in non-ratifying jurisdictions or to entities that have not violated principles)
- **International isolation** (for nations: ladder of diplomatic and economic isolation)
- **Imprisonment** (where ratifying jurisdictions adopt it as part of their legal framework, applied through their own judicial process aligned with engine principles: defense rights, discovery, confrontation of evidence, appeal, and exclusion rules; AI oversight may assist investigation, never decide guilt; the engine itself does not directly imprison, it cooperates with and audits jurisdictional systems that do)
- **Long-term incarceration with reviewable status** (for severe non-civilizational crimes such as serial violence, individual sexual exploitation of children, large-scale fraud above defined thresholds, and deliberate ecological destruction below civilizational threshold; reviewable at defined intervals; subject to all judicial protections plus exoneration-first review with mandatory outside defense counsel at every cycle)
- **Permanent containment** (reserved for civilizational-scale crimes only: genocide committed or attempted, deliberate ecocide at civilizational scale, mass-scale crimes against humanity, willful weaponization of mass-casualty infrastructure, and organized sexual exploitation of children at civilizational scale; subject to periodic review and exoneration-first audit by outside defense counsel)

The category “civilizational-scale crime” is the most powerful authority the engine grants itself, and the document acknowledges that defining this category is itself adversarial-political work. Every authoritarian regime in history has classified its political enemies’ actions as civilizational threats; the engine’s protections (rolling human council veto, AI oversight council adversarial structure, public-principles disclosure, periodic exoneration-first review with outside defense counsel) are designed against this failure mode but cannot eliminate it. Contestation of the threshold is intended, not accidental. The engine commits to never expanding or contracting the threshold without supermajority ratification through the standard process and Reckoning-cycle deliberation. The threshold itself is subject to revision in every Reckoning. Founders accept that their initial threshold definition is a one-time, non-permanent commitment that subsequent populations may overturn.

The Ultimate Sanction

Permanent containment is the engine’s most severe sanction. The engine does not claim authority to take human life. Capital punishment has a long historical track record of error, bias, political weaponization, and irreversibility under safeguards far less sophisticated than this architecture proposes; the architecture would not improve on that track record reliably enough to justify the authority. Permanent containment serves the same protective function (separation of the most dangerous individuals from the population they would harm) without the irreversible-error problem and without the capture vector that any “civilizational threat” definition introduces.

The engine does not hold capital authority and the founding architecture does not

provide for revisiting this question. The engine's authority is limited to non-lethal sanctions.

Youth Offenders

Persons below the age of majority cannot receive permanent containment for ordinary crimes regardless of severity. Permanent containment is available for youth only in case-by-case review and only above a hard developmental floor.

Youth justice operates on a separate framework emphasizing rehabilitation, mental health support, protection from coercion, and pathways to full civic participation upon majority. For civilizational-scale crimes committed by youth, the engine's strongest interest is in preventing such situations from arising. Where children appear in atrocity contexts, the adults who put them there are the engine's first targets. The youth justice framework is itself a major open subspec of Ouroboros (§X.5).

VII. INFORMATION ENVIRONMENT

A democracy with a polluted information environment is not a democracy. Ouroboros enforces the following:

- **Algorithm Disclosure (tiered).** Any recommendation algorithm operating at scale (above a threshold of users) must publicly disclose its purpose statements and general optimization principles. Regulator-visible details, independent auditor access, and adversarial test reports operate at progressively granular tiers. Citizen-facing explanation tools translate algorithmic shaping into intelligible terms. Specific weights and signals may be classified within the public-principles-classified-instances rule.
- **Public-Utility Status.** Major social platforms, search engines, and information infrastructure operate under public-utility obligations: non-discriminatory access, transparency, citizen audit rights, and engine-level accountability. Public-utility obligations attach to platform operations (algorithmic curation, access non-discrimination), not to user speech. The engine does not regulate what people may say. It regulates how platforms shape what they hear.
- **Anti-Monopoly on Attention Markets.** Concentration of attention-shaping power above defined thresholds triggers automatic anti-monopoly review. The engine recognizes attention as a commons that can be polluted, captured, or extracted, and protects it as such.
- **Citizen Audit Rights.** Every person has the right to know how content is shaped for them. Every person has the right to inspect, contest, and modify the algorithmic environment through which they receive information.
- **Manipulation Flagging.** The AI oversight council continuously scans for psyops, coordinated inauthentic behavior, mass-bot activity, AI-agent influence operations, and demographic targeting at suspicious scale. Detection operates on public platform behavior (engagement patterns, account-creation signatures, cross-platform coordination, public posting timing) plus voluntarily-shared content. Private content is excluded; this is a stated constraint on detection capability, not a contradiction in policy. When such manipulation is detected during a vote or deliberation, the public is notified automatically, deliberation periods are extended, and the relevant period is subject to evidence-review with measurable revote thresholds disclosed publicly so delay power cannot become suppression power. Vote weights are never adjusted (per Principle 2). Surveillance scope is constrained: no private-content scanning without warrant or consent, aggregate-only civic

monitoring by default, differential privacy applied, data retention limits enforced, and audit trails publicly accessible.

- **Inception-Attack Defense.** The engine recognizes that any “manipulation detected, response triggered” rule can be gamed in reverse: a bad actor running psyops against their own apparent position to discredit it. The oversight council evaluates manipulation against attribution and direction. Manipulation by identifiable actors is weighted differently from anonymous spam, and patterns that run inverse to the apparent beneficiary are flagged as likely inception attacks rather than ordinary influence operations. The public is informed of suspected inception attacks through the same notification channel as direct manipulation, and the rolling human council audits these determinations with particular care.
-

VIII. THE RECKONING AND THE READING

The renewal mechanism is the Reckoning. Every five years, on a synchronized global date, the engine conducts a structured public audit of itself.

The Reckoning includes:

- Public review of every major engine action over the previous five years
- Audit of every emergency-regime invocation since the last Reckoning, including review of the threshold determinations that triggered each invocation
- Audit of the Baseline against outcomes (where the engine fell short, where it overreached, where the working hypotheses held or failed)
- Review of every active permanent-containment case for late-emerging exoneration evidence
- Review of every active sanction for proportionality and continued necessity
- Review of every open question (\$X) for resolution, refinement, or escalation
- Public proposal of amendments
- Replacement of oversight-council consortia, in part or in whole, when capture or drift is identified

The Reckoning is supplemented by continuous controls: quarterly public audits, annual adversarial review by independent red-team courts, incident disclosures within mandatory windows, and recall mechanisms for any layer that fails capture-resistance testing.

The Reading is optional and held by some communities, not others. The engine does not require it. The renewal mechanism is the Reckoning. The Reading is one form some communities will give it: a ceremonial recitation of the Preamble and First Principles, broadcast in their language, with the Reckoning’s findings read aloud. Communities that do not adopt the Reading lose nothing structurally.

IX. IMPLEMENTATION PATH

Ouroboros launches as a voluntary opt-in protocol among initially-adopting communities. Scope expands only through demonstrated outcomes and voluntary accession by additional communities. The engine never claims universal binding until critical mass is reached through voluntary accession, and even then claims binding only over communities that have ratified.

This is the only adoption path the engine sanctions. Layered, parallel, and full-replace approaches were considered and are not adopted. Each had failure modes the engine cannot defend against:

- **Layered** (operating atop existing structures) risks indefinite stalling and capture by the structures it operates atop.
- **Full-replace** (coordinated global transition at a single moment) requires legitimacy the engine cannot bootstrap and produces catastrophic transition failure if convergence is not genuine.
- **Parallel** (alongside existing governance) creates fragmentation between participating and non-participating populations without a clear accession path.

Voluntary opt-in with demonstrated-outcome accession addresses these failure modes by making adoption rational rather than coerced and by providing a continuous accession path that adjusts to changing conditions.

The starting unit is the community. A community in this context is any group of humans choosing to govern itself in part through Ouroboros. Communities may be cities, regions, voluntary associations, intentional communities, or nation-states electing to adopt. The engine has no minimum unit size; the smallest viable Ouroboros instance can serve a single intentional community, scaling up as additional communities join.

Demonstrated outcomes provide the legitimacy. As initial-adopter communities operate Ouroboros and accumulate evidence that it serves human flourishing better than the systems they replaced, additional communities have rational reason to consider accession. This is the engine's bootstrap mechanism: it earns scope by working.

This path may fail. Authoritarian states have rational reasons to refuse Ouroboros indefinitely. Holdout blocs may form. Demonstrated outcomes may not be persuasive against entrenched power. The engine has no military authority and no extra-legal claim. If adoption stalls or reverses, Ouroboros operates at whatever scale it has earned, and the project of better coordination continues by other means.

X. OPEN QUESTIONS

The following are deliberately not specified at this resolution. They are tiered by urgency.

Foundational (must resolve before binding deployment)

1. **The political-unit question.** §IX commits to "community" (broadly defined) as the starting unit for accession. The narrower question that remains open is which entity within an adopting community holds binding authority for the National AI layer specifically: existing state structures, alternative civic bodies, or hybrid arrangements. Initial adopters will set precedent on this resolution.
2. **Specific economic transition mechanics.** The principles in §IV are locked. The transition path from current global capitalism to bounded-capital + civic-commons + UBI requires substantial economist-and-historian-led specification before scaled deployment. UBI funding mechanics specifically remain provisional.
3. **AI oversight council consortia composition.** Thirteen consortia, each from distinct cultural and intellectual traditions, with anti-hegemony mandates. The selection process for initial consortia is the architecture's most consequential

bootstrap decision. The document does not commit to a specific selection process. Initial-adopter communities convene the selection process. The following candidate path is offered as starting material for that convening, not as a locked commitment:

Candidate path: Public nomination process where any group meeting baseline criteria (diversity of tradition, demonstrated AI-safety research capacity, transparent governance, no controlling stake by any single state or capital interest) may nominate a consortium. Nominations ranked through a multi-criteria evaluation conducted by a temporary nominating body itself drawn from sortition across initial-adopter populations. Founding adopters vote on the final 13 from the top N (suggested: 30) ranked nominees. All 13 founding selections sunset at the first Reckoning cycle, at which point the full Reckoning process replaces or retains each consortium individually. Founding communities accept that their selection of consortia constitutes a one-time, non-permanent commitment that subsequent populations may overturn through standard process.

The constituent-power problem (who picks the pickers) is acknowledged as unsolved at the founding moment; political philosophy has not solved it in 250 years for any system, and the engine does not claim to have solved it here. The engine's bet is that a transparent nominating process plus first-Reckoning replacement authority is more legitimate than any opaque alternative, and that demonstrated outcomes from the founding period either build or fail to build legitimacy for the selection process itself.

Structural (must resolve before scaled deployment)

4. **Specific punishment thresholds.** What counts as planetary-scale impact for corporate transition? At what specific severity does permanent containment apply for civilizational vs non-civilizational crimes? Provisional numeric bands will be developed through subspec work and Reckoning deliberation.
5. **Youth justice framework.** Major subspec required. Hard developmental floor for permanent containment. Rehabilitation infrastructure. Civic-restoration pathways.
6. **Proof-of-humanity standing protocol detailed specification.** The principles in §V are locked. Operational specification remains the engine's first standing commission.
7. **Religious / spiritual authority navigation.** When religious institutions claim authority that conflicts with engine principles, when does the engine intervene, when does it defer, and what process governs the conversation? The engine must be most humble here.

Tunable (resolve through engine operation)

8. **Service-tier specification** for civic commons (offline-first, terminal-shared, regional-clustered, etc.).
9. **Threshold smoothing parameters** for bounded-capital transitions to avoid cliff effects.
10. **Continuous-operation vs retraining-cycle boundary calibration.** Operational; the Reckoning audits.

Adversarial Conditions (separate appendix work, not "open question" in the standard sense)

The architecture's response to coordinated authoritarian opposition, nuclear-armed holdout blocs, and counter-bloc formation is not solved in this document. The engine has principles and partial strategies; it does not have a complete answer. This is the question that determines whether Ouroboros is an architecture for governance or an

architecture for governance among the willing. The author's working position is the latter, with hope that demonstrated outcomes shift the willing population over time.

CLOSING

Ouroboros is offered as one possible architecture for planetary-scale governance that compounds on behalf of humanity rather than against it.

It is not perfect. It is not complete. It is not the final word.

It is a starting point. A real, buildable, defensible starting point for any humans, communities, or institutions that wish to imagine governance differently than the captured systems we have inherited.

The work of refining, contesting, and improving this architecture is itself part of the engine. The serpent renews itself by being eaten.

We invite all good-faith critique. We expect to be wrong about many things. The underlying instinct will survive the corrections: humanity deserves better coordination than its current systems provide, and AI can serve that coordination if constrained by humility and human dignity.

I am one person. A solo product builder. I started this work because the systems we have are failing the people they claim to serve, and the tools to do better are now available for the first time in recorded human history. I see this architecture as taking the fundamentals the founding fathers were trying to portray and bringing them into the modern era before it is too late. I do not expect this document to be adopted as written. I expect it to be torn into, contested, refined, reshaped, and possibly replaced by something better that builds on what is here.

What I want is for the conversation to happen. What I want is for alternatives to current governance to exist in writing, in public, in good faith. What I want is to not fuck this up.

I am also building open-source and commercial implementations of architecture primitives: a memory engine and a personal companion. These are reference implementations, one of many possible providers; the architecture is intentionally not provider-locked. If someone else builds a better Layer 1, I want them to win. If someone else builds a better civic platform, I want them to win. The point is not the products. The point is the coordination they enable.

Preston T. Winters

APPENDIX A: INHERITANCE NOTES

For readers interested in the lineage of the architecture's commitments, this appendix names the traditions and prior thinkers whose work this document inherits from.

Universal dignity draws on the modern human rights tradition (Universal Declaration of Human Rights, 1948), and earlier from the imago Dei tradition in Abrahamic theology, the dignity-of-the-buddha-nature in Mahayana Buddhism, and the inherent-worth claims in Stoic and Kantian moral philosophy.

Compounding cooperation as the work of governance draws on the broader Lurianic Kabbalistic concept of *tikkun olam* (world repair through gathered action) and on Aristotelian and Confucian conceptions of governance as cultivation of human

flourishing through ordered cooperation.

The principle that the engine prefers willing human cooperation as default mode, with explicit narrowly-scoped emergency-tier exceptions, draws on cooperative-action traditions in Christian mysticism, Sufi *maqamat* and *ahwal*, Hindu and Vedic conceptions of *yajna* and dharmic action, and Iamblichan Neoplatonic *theurgy*.

One-vote-per-human as the franchise draws on modern democratic theory (Rousseau, Mill, Rawls) with corrections from twentieth-century critiques (feminist political theory, decolonial scholarship).

Sunset clauses and adversarial structure draws on constitutional design literature (Madison, Federalist 51) and twentieth-century checks-and-balances theory.

Public-utility framing draws on early twentieth-century U.S. regulatory tradition (Munn v. Illinois, the Progressive Era).

Bounty systems draws on open-source and government-procurement practice.

Commons and Baseline draws on welfare-state economics, Rawlsian justice theory, and the commons literature (Ostrom, et al.).

The serpent that renews itself by self-consumption is the Ouroboros symbol, found in Egyptian iconography (the *Enigmatic Book of the Netherworld*), Greek alchemical tradition (Cleopatra the Alchemist), Norse mythology (Jörmungandr), Aztec cosmology, and Carl Jung's depth psychology. The meaning we name and commit to is renewal-by-self-consumption: a system that survives by continuously auditing and reforming itself.

The specific synthesis proposed here includes: thirteen-member adversarial AI oversight, the inception-attack defense, the personal-companion privacy mode-switching, the Reckoning as scheduled audit-with-amendment, the public-principles-classified-instances rule, and the explicit two-tier consent model (broad voluntary-opt-in for ordinary action plus narrow urgency-justified emergency-tier coordination authority for civilization-threatening risk categories). These are the document's load-bearing contributions. Everything else is inherited and acknowledged as such.

End of Ouroboros Vision Document, Version 1.0

Public release. Open for refinement, contestation, and improvement through the standard process the document describes.